

Implementing a Round-Robin Scheduler Using pthread library

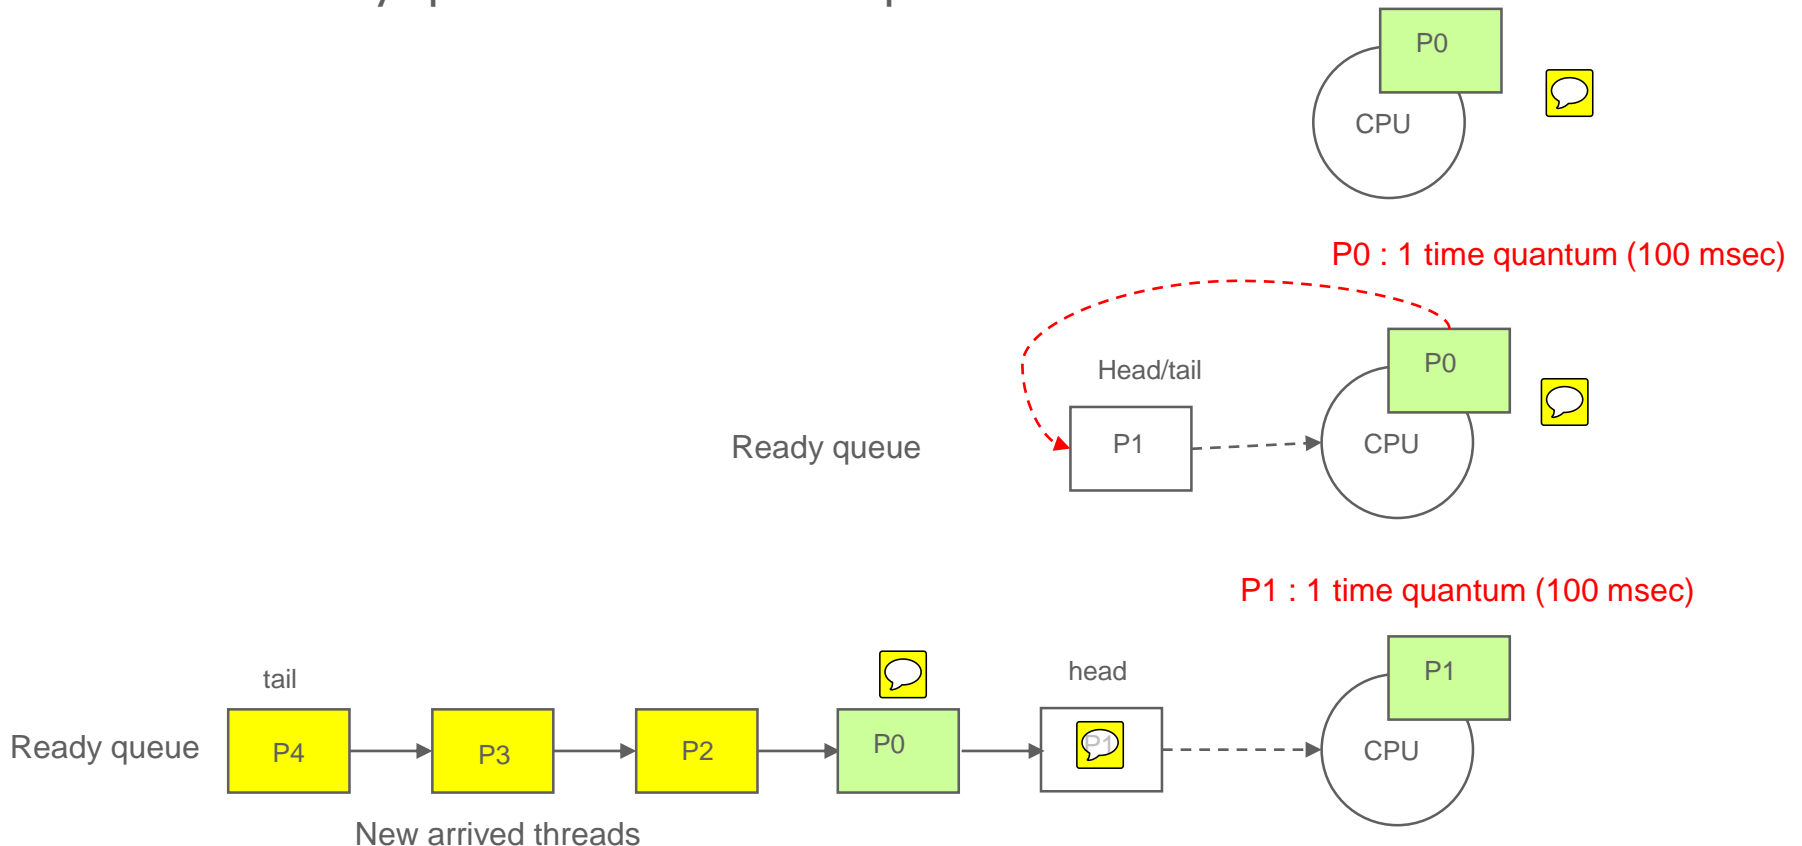


Woo Hyun Ahn (whahn@kw.ac.kr)

Round Robin (RR) Scheduling


- Scheduling methodology

- Each process with **the same priority** gets a small unit of CPU time (*time quantum or time slice*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue. Then the next process in FIFO is executed.




Thread Control Block (TCB)

■ Thread Control Block

```
typedef _thread {  
    ThreadStatus      status;  
    int exitCode;  
    pthread_t         tid;  
    pthread_cond_t    readyCond;  
 BOOL               bReady;  
    pthread_mutex_t   readyMutex;  
    Thread*           pNext;  
    Thread*           pPrev;  
} Thread;
```

■ Thread status

```
Enum {  
    THREAD_STATUS_RUN = 0,  
    THREAD_STATUS_READY = 1,  
    THREAD_STATUS_SLEEP = 2,  
    THREAD_STATUS_ZOMBIE = 3   
} ThreadStatus;
```

Thread Control Block (TCB)

- System data structure including thread-specific information.
- TCB contains everything a kernel needs to know about a particular thread.
- Thread status, priority, name, parent/child task information, etc.

Creating a thread with thread



A thread is created with

```
int thread_create(  
    thread_t *thread,  
    const thread_attr_t *attr,  
    void *(*start_routine)(void *),  
    void *arg);
```

- The creating thread must provide a location for storage of the thread id.
- The third parameter is just the name of the function for the thread to run.
- The last parameter is a pointer to the arguments.

The Thread ID



`thread_t thread_self(void)`

- Each thread has an id of type `thread_t`.
 - On most systems this is just an integer (like a process ID)
 - But it does not have to be

thread_join and thread_exit

```
int thread_join(thread_t thread, void** retval);
```

- thread_join() is a blocking call on threads
- It indicates that the caller wishes to block until the thread being joined exits.

- ```
int thread_exit(void* retval);
```

  - Should be called before thread is terminated.

# Thread\_suspend(thread\_t tid)



- Suspends a thread.
- SYNOPSIS
  - `int thread_suspend( thread_t tid);`
- Parameters
  - *tid*
    - [in] thread ID of a thread to suspend.
- Return Values
  - If the function succeeds, the return value is 0; otherwise, it is (-1).


# Thread\_resume(thread\_t tid)

- Resume a specific suspended task.
- **SYNOPSIS**
  - `int thread_resume( thread_t tid);`
- **Parameters**
  - *tid*
    - [in] thread ID of a thread to resume.
- **Return Values**
  - If the function succeeds, the return value is 0; otherwise, it is (-1).




# Ready Queue & Task Waiting Queue

Ready Queue

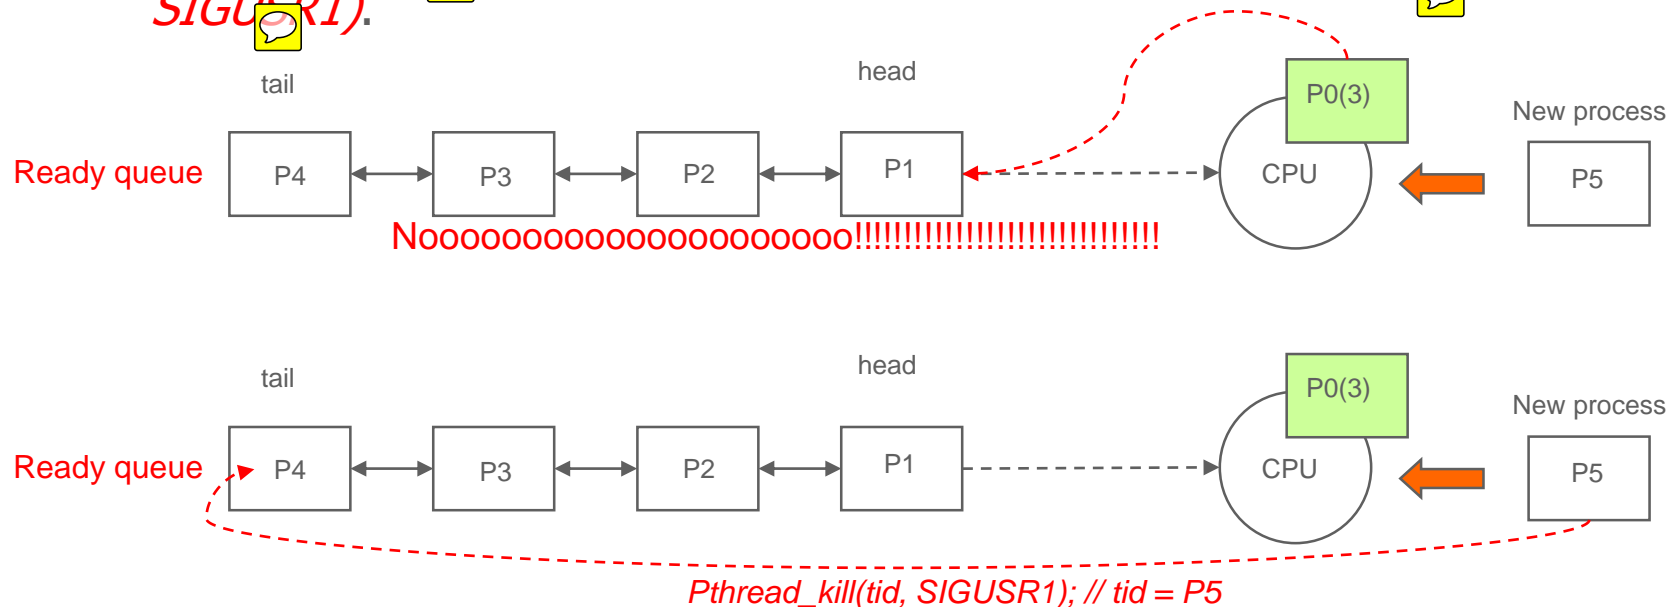


Waiting Queue

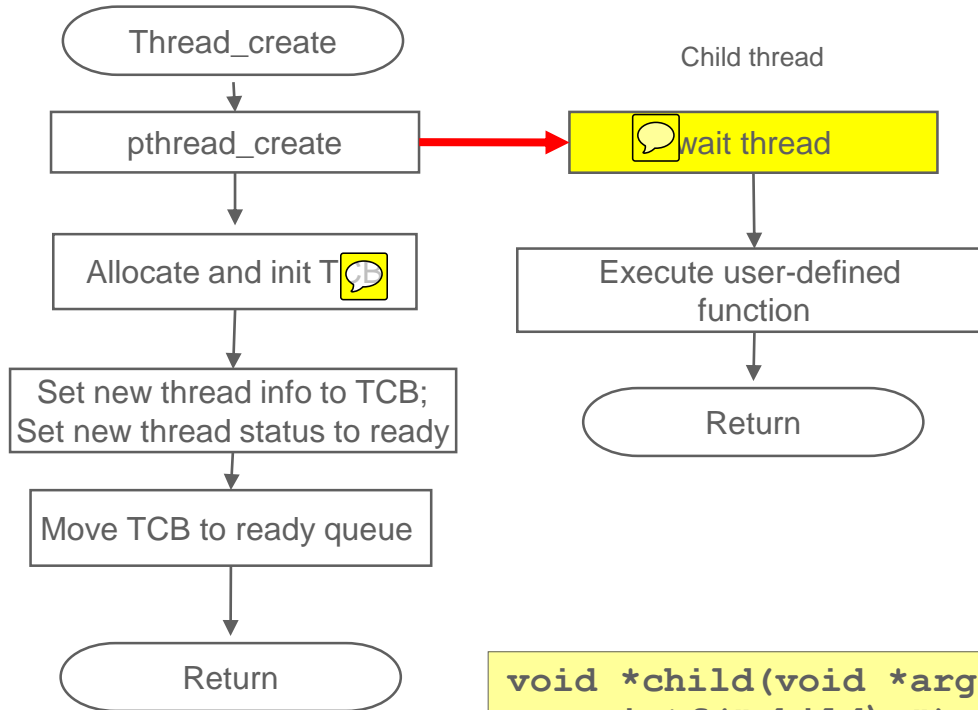


# Ready Queue & Sleep Queue (Cont'd)

- A newly created thread is not run soon, but first moved at the tail of the ready queue.
  - Threads in the ready queue should be executed in the round-robin manner.
  - The created thread should be waited until the scheduler sequentially execute other threads that are placed in the ready queue.
  - When a thread is created, the thread is waited by *\_\_thread\_wait\_handler*, and the thread is executed by *pthread\_kill(tid, SIGUSR1)*.



# thread\_create



- Child thread is suspended when it is created.
  - Otherwise, the scheduler cannot control the child thread.
- How can the child execute the waiting code?
  - Use a wrapper function that includes the user-defined function when calling pthread\_create function.


```
void *child(void *arg) {
 printf("child\n");
 return NULL;
}
int main(int argc, char *argv[])
{
 pthread_t c;
 pthread_create(&c, NULL, child, NULL);
 pthread_join(c);
 return 0;
}
```

# Wrapper Function

- Thread\_create creates a child calling the wrapper function.

```
void *child(void *arg) {
 printf("child\n");
 return NULL;
}

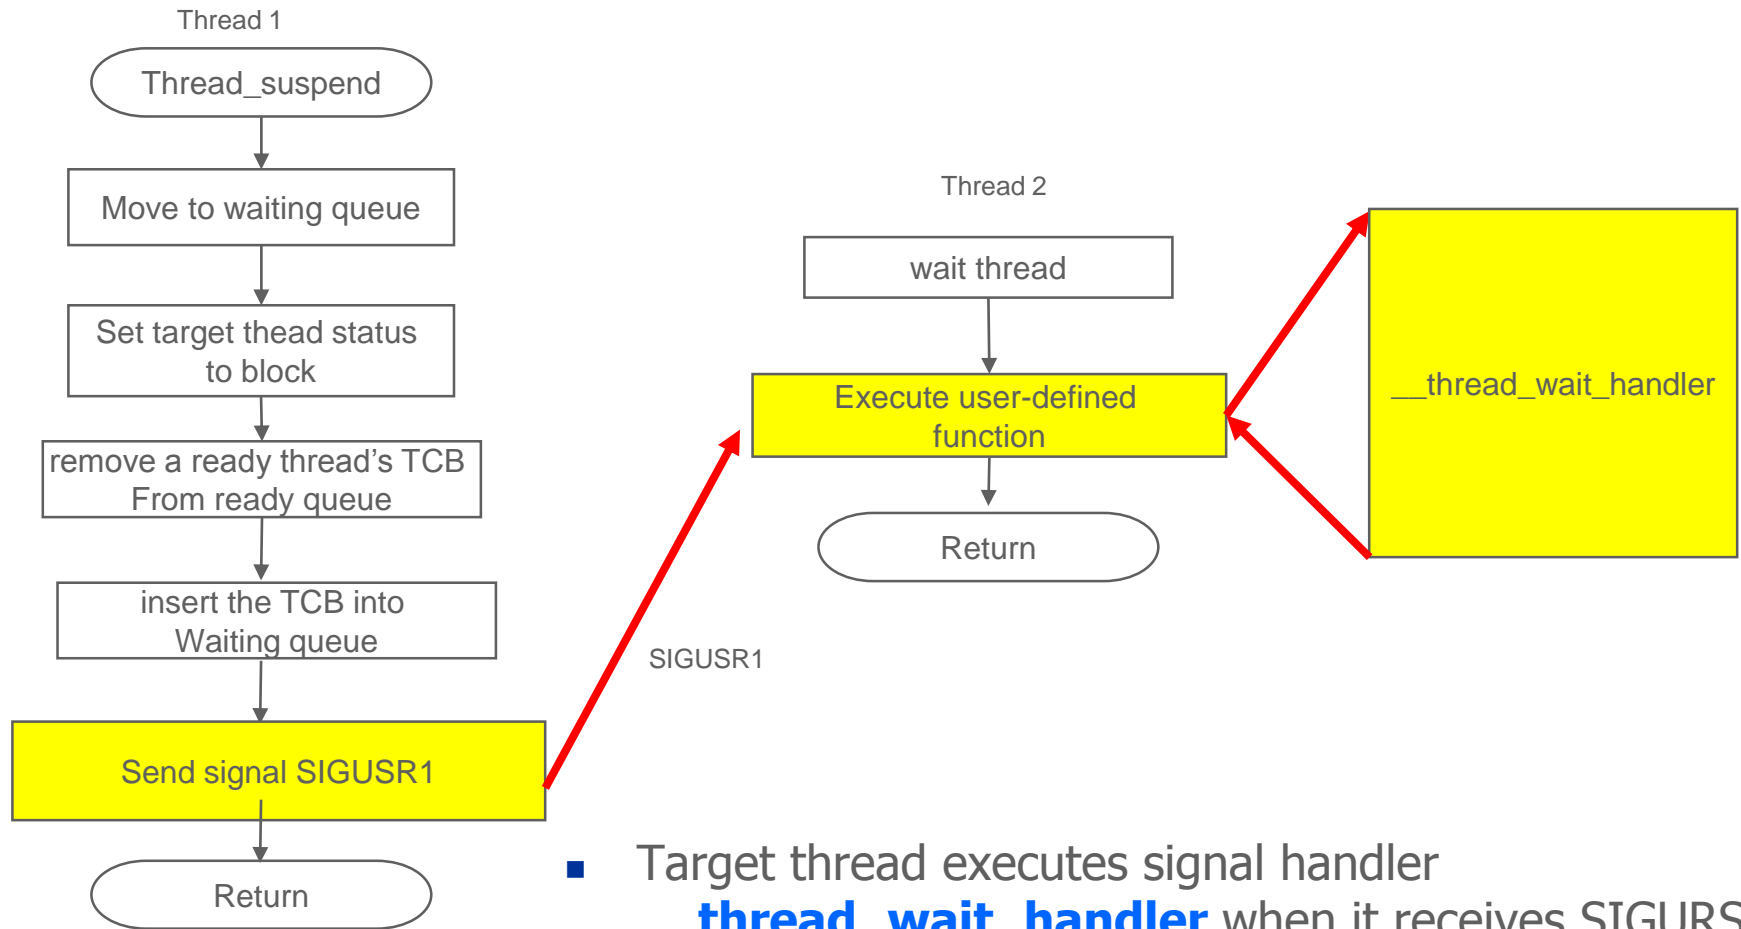
void* __wrapperFunc(void* arg)
{
 void* ret;
 WrapperArg* pArg = (WrapperArg*)arg;

 void* funcPtr = pArg.funcptr;
 void* funcArg = pArg.funcArg;
 
 ret = (*funcPtr)(funcArg);
 return ret;
}
```

```
Typedef __wrapperArg {
 void __ (*funcPtr)(void*);
 void* funcArg;
} WrapperArg;

int main(int argc, char *argv[])
{
 pthread_t c;
 WrapperArg wrapperArg;
 wrapperArg.funcptr = child;
 wrapperArg.funcArg = 10;
 Pthread_create(&c, NULL,
 __wrapperFunc, &wrapperArg);
 pthread_join(&c);
 return 0;
}
```

# thread\_suspend



- Target thread executes signal handler `___thread_wait_handler` when it receives `SIGUSR1`
- A thread is suspend in a signal handler `___thread_wait_handler` when it receives signal `SIGUSR1`



# Thread Wait using Signal Handler

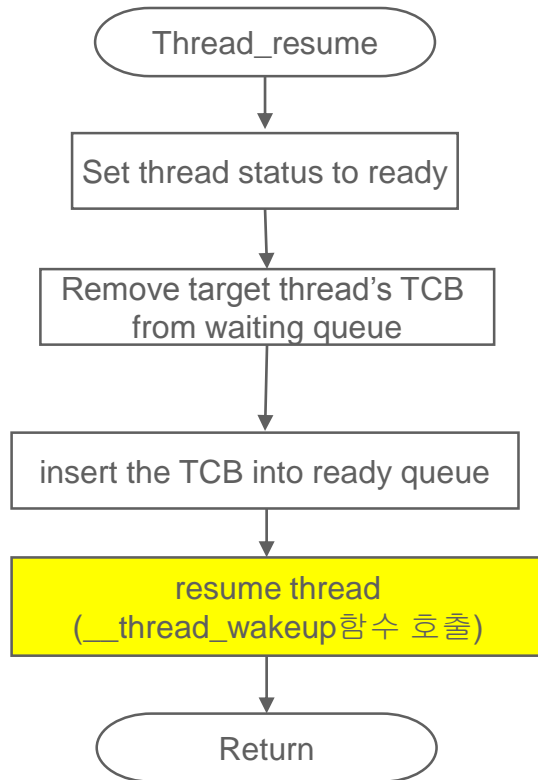


```
void __thread_wait_handler(int signo)
{
 Thread* pTh;

 pTh = __getThread(pthread_self());
 pthread_mutex_lock(&(pTh->readyMutex));
 while (pTh->brunable == FALSE)
 pthread_cond_wait(&(pTh->readyCond), &(pTh->readyMutex));
 pthread_mutex_unlock(&(pTh->readyMutex));
}
```

- Signal handler for thread waiting
  - Thread\_suspend sends signal SIGUSR1 to the target thread.
  - \_\_thread\_wait\_handler is invoked to sleep the target thread.

# thread\_create, thread\_suspend, thread\_resume



# \_\_thread\_wakeup

- \_\_thread\_wakeup
  - Signals on the condition variable to wake up the target thread that is blocked on the condition variable in \_\_thread\_wait\_handler()

```
void __thread_wakeup(Thread* pTh)
{
 pthread_mutex_lock(&(pTh->readyMutex));
 pTh->bRunnable = TRUE;
 pthread_cond_signal(&(pTh->readyCond));
 pthread_mutex_unlock(&(pTh->readyMutex));
}
```

# Thread join

Example 1

```
int done = 0;

void *child(void *arg) {
 printf("child\n");
 done = 1;
 return NULL;
}

int main(int argc, char *argv[])
{
 pthread_t c;
 printf("parent: begin\n");
 pthread_create(&c, NULL,
 child, NULL);
 while (done == 0); // spin
 printf("parent: end\n");
 return 0;
}
```

Example 2

```
int done = 0;
pthread_mutex_t m =
 PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t c =
 PTHREAD_COND_INITIALIZER;
void *child(void *arg) {
 printf("child\n");
 pthread_mutex_lock(&m);
 done = 1;
 pthread_cond_signal(&c);
 pthread_mutex_unlock(&m);
 return NULL;
}

int main(int argc, char *argv[])
{
 pthread_t c;
 printf("parent: begin\n");
 pthread_mutex_lock(&m);
 while (done == 0)
 pthread_cond_wait(&c, &m);
 pthread_mutex_unlock(&m);
 printf("parent: end\n");
 return 0;
}
```

# Example: Thread join

Example 2

```
int done = 0;
pthread_mutex_t m =
 PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t c =
 PTHREAD_COND_INITIALIZER;
void *child(void *arg) {
 printf("child\n"); mythread_exit()
 pthread_mutex_lock(&m);
 done = 1;
 pthread_cond_signal(&c);
 pthread_mutex_unlock(&m);
 return NULL;
}
int main(int argc, char *argv[])
{
 pthread_t c; mythread_join()
 printf("parent: begin\n");
 pthread_mutex_lock(&m);
 while (done == 0)
 pthread_cond_wait(&c, &m);
 pthread_mutex_unlock(&m);
 printf("parent: end\n");
 return 0;
}
```

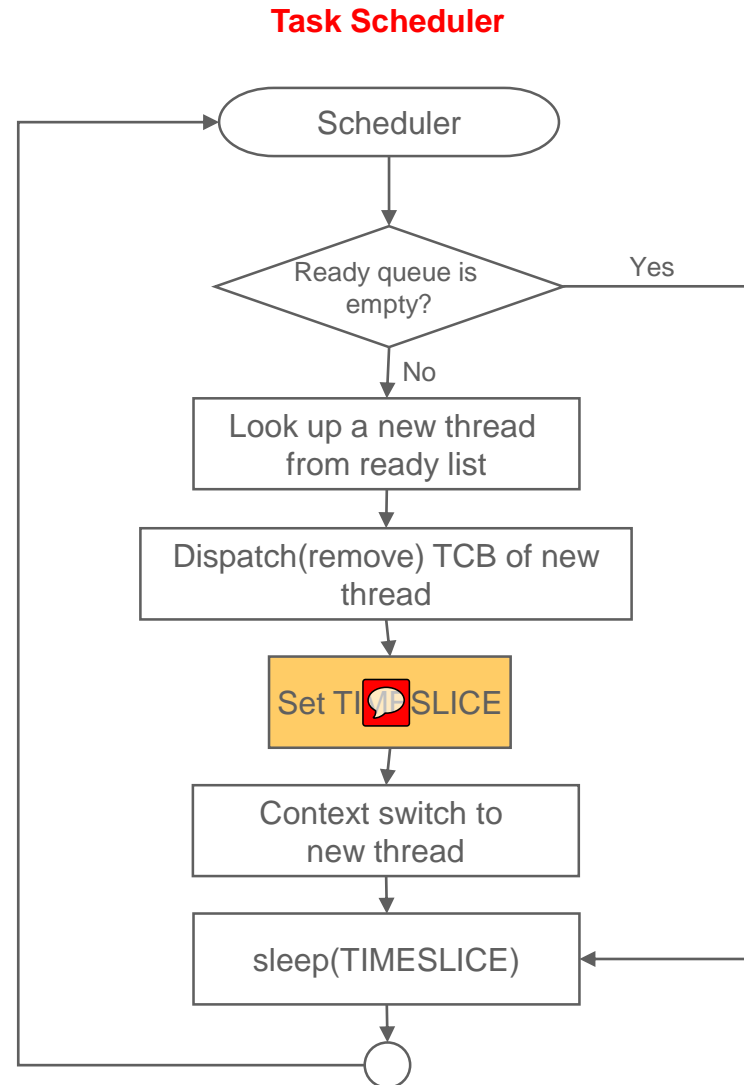
Example 3

```
int done = 0;
pthread_mutex_t m =
 PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t c =
 PTHREAD_COND_INITIALIZER;
void *child(void *arg) {
 printf("child\n");
 mythread_exit();
 return NULL;
}
int main(int argc, char *argv[])
{
 pthread_t c;
 printf("parent: begin\n");
 mythread_join();
 printf("parent: end\n");
 return 0;
}
```



# Thread Scheduler

- Task scheduler is a background thread that is not seen to applications.
  - The initialization thread (or main thread) is the scheduler thread
- Time slice
  - `#define TIMESLICE (2)`
- `void RunScheduler(void)`
  - Run thread scheduler.
  - Implemented by `sleep()`



# Context Switching

## ■ Procedures

- Stop the current running thread: *pthread\_kill(curtid, SIGUSR1)*
- Execute a given target thread that is intended to be executed (i.e., scheduled, or dispatched) in the next order:  
*\_\_thread\_wakeup(pNewThread)*
- pNewThread can be obtained by tid of the target thread to be executed.

## ■ Interface

- `void _ContextSwitch(Thread pCurThread, Thread* pNewThread)`


# System Initialization

- System initialization routine creates the following:
  - Scheduling queues (ready queue and waiting queue)
  - Thread scheduler.

## Main.c

```
void main()
{
 thread_id tid;
 int arg;
 Init();

 thread_create(&tid, NULL, AppTask, &arg);

 RunScheduler(); 
}
```

## Init.c

```
void Init(void)
{
 // Create ready queue and waiting queue
 ...
 // initialize thread scheduler
}
```

## App.c

```
Void* AppTask(void* param)
{
 TestCase();
 return NULL;
}
```

# Testcase

Testcase.c

```
void* foo1(void* arg)
{

 while(1);
}

void foo2(void* arg)
{

 while(1);
}
...

void Testcase(void)
{
 int tid1, tid2, tid3;
 int arg1, arg2, arg3;

 thread_create(&tid1, NULL, foo1, &arg1);
 thread_create(&tid2, NULL, foo2, &arg2);
 thread_create(&tid3, NULL, foo3, &arg3);
 ...
 thread_suspend(tid1);
 thread_suspend(tid2);
 ...
 thread_resume(tid1);
 while(1);
}
```

# Building

